

Exploratory Data Analysis

- Input Data is thousands of satellite images taken of country, city, and other landscapes.
- Goal: Predict which pixels are roads, and which pixels are not.

Sample Image




Paired Mask




Sample Image for ANN




Average Road Pixel



Average Non-Road Pixel




Ground Truth vs Prediction

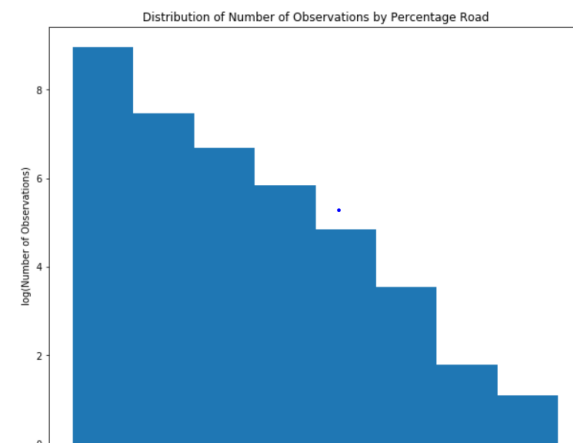


- We observe that the average road pixel is slightly lighter than the non-road pixel, but there isn't an easily identifiable difference
- We attempted a simple ANN model which predicted on 3x3 patches

Relative Frequency of Road Pixels



Distribution of Proportion of Roads Pixels



- There are more road pixels on the right side of the images
- There is an exponential dropoff in frequency as the percentage of road pixels increases. Many more examples with very few road pixels!


Preprocessing Steps

- Mean and Variance Scaling:**
 - Standardizing training data using a Z-score transformation at each pixel
 - Build model on variations instead of raw images to ease prediction

Original




After Z-Score Standardization



- Data Augmentation:**
 - Model will be able to generalize to unseen images better
 - Randomly apply to input images


Brightness Transform



Rotation




Horizontal Flipping



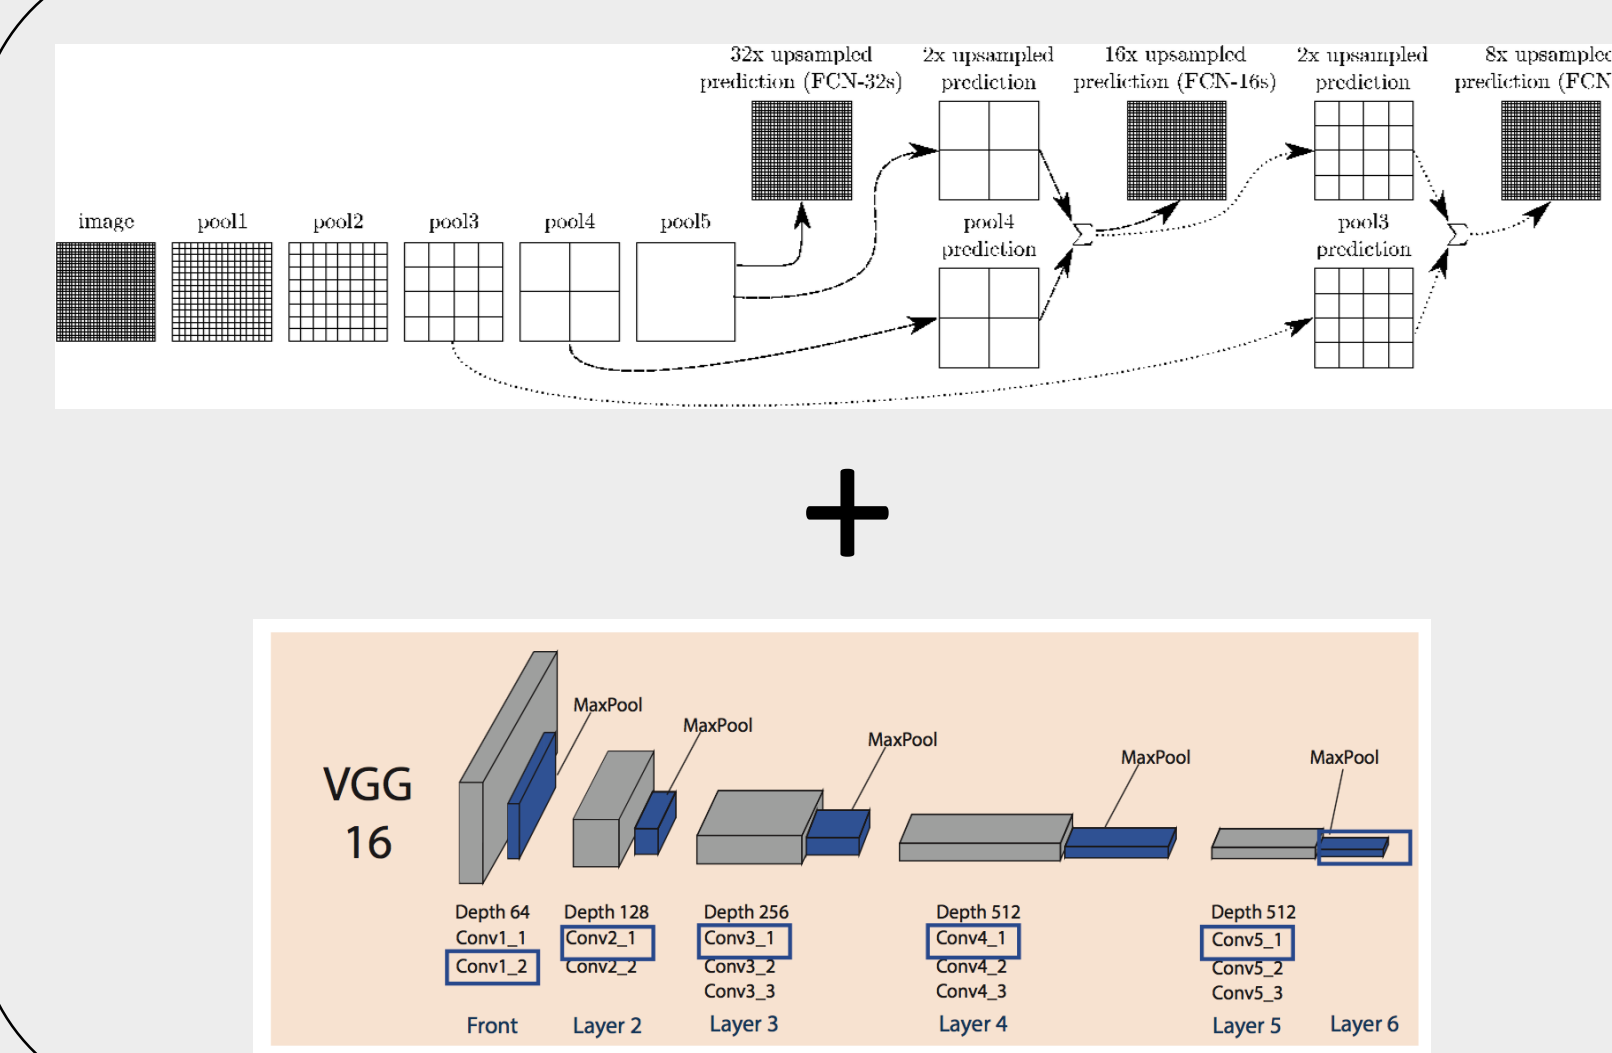
Vertical Flipping



Final Model




X 5



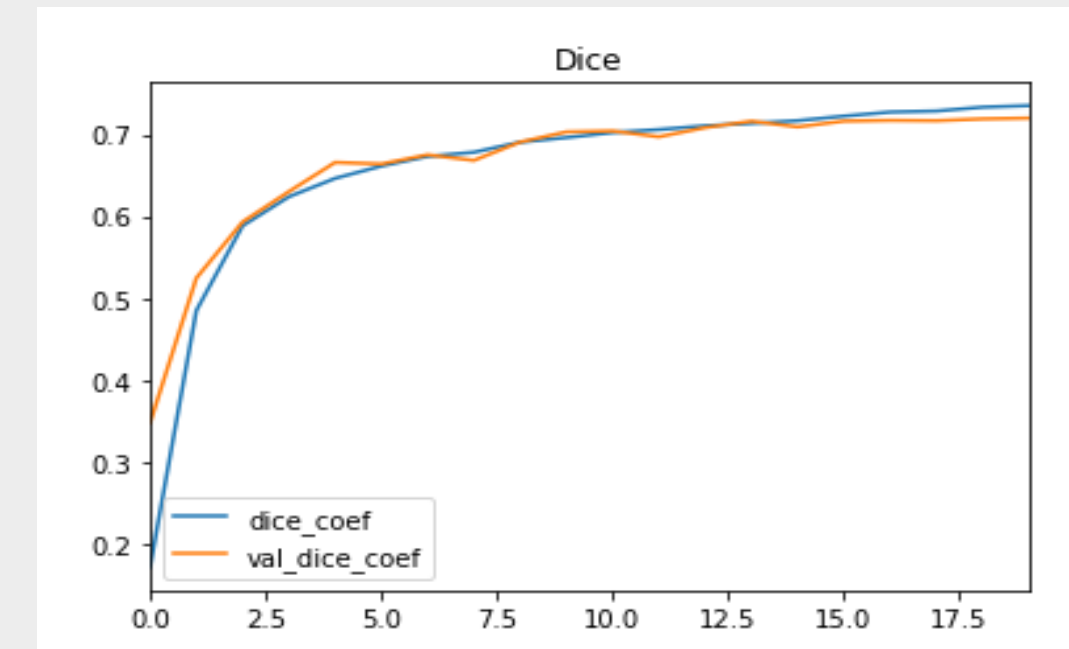
- Our final model was an ensemble of 6 convolutional neural networks consisting of 3 standard U-Nets with dropout layers, 2 modified U-Nets with batch normalization instead of dropout layers, and an FCN-8 trained starting from pre-trained VGG-16 weights
- Ensembling boosted our accuracy by around 3%
- The 3 standard U-Nets consisted of the following: 40 epochs/no augmentation, 50 epochs/no augmentation, 50 epochs/augmentation
- The 2 modified U-Nets consisted of the following: 30 epochs/augmentation, 40 epochs/augmentation

Hyperparameter Selection

Learning Rate = 1e-4



Learning Rate = 1e-3



Learning Rate = 1e-2



- The learning rate of 1e-3 was used to train each U-Net as well as the FCN-8
- The maximum batch size which fit into memory was always used as a larger batch size typically allows for a more accurate gradient calculation (32 for the U-Nets and 1 for the FCN-8)
- The Adam optimizer was used as it is generally agreed to be the most efficient
- The training data was split as such: 90% training, 10% validation (large amount of training data is effective, and validation size still reasonable)

Post-Processing Steps

Conditional Random Fields (CRF)

$$k(\mathbf{f}_i, \mathbf{f}_j) = w^{(1)} \exp \left(-\frac{|\mathbf{p}_i - \mathbf{p}_j|^2}{2\sigma_p^2} - \frac{|\mathbf{I}_i - \mathbf{I}_j|^2}{2\sigma_I^2} \right) + w^{(2)} \exp \left(-\frac{|\mathbf{p}_i - \mathbf{p}_j|^2}{2\sigma_p^2} \right)$$


Tutorial CRF Results



Our Results



OpenCV Image Processing





Ground Truth



Feeding Predicted Masks into U-Net

- The predicted masks for all the training examples generated by our top model were used as new training examples paired with the original training masks as labels
- The default U-Net was then trained on this new dataset
- Increased accuracy by around 2%

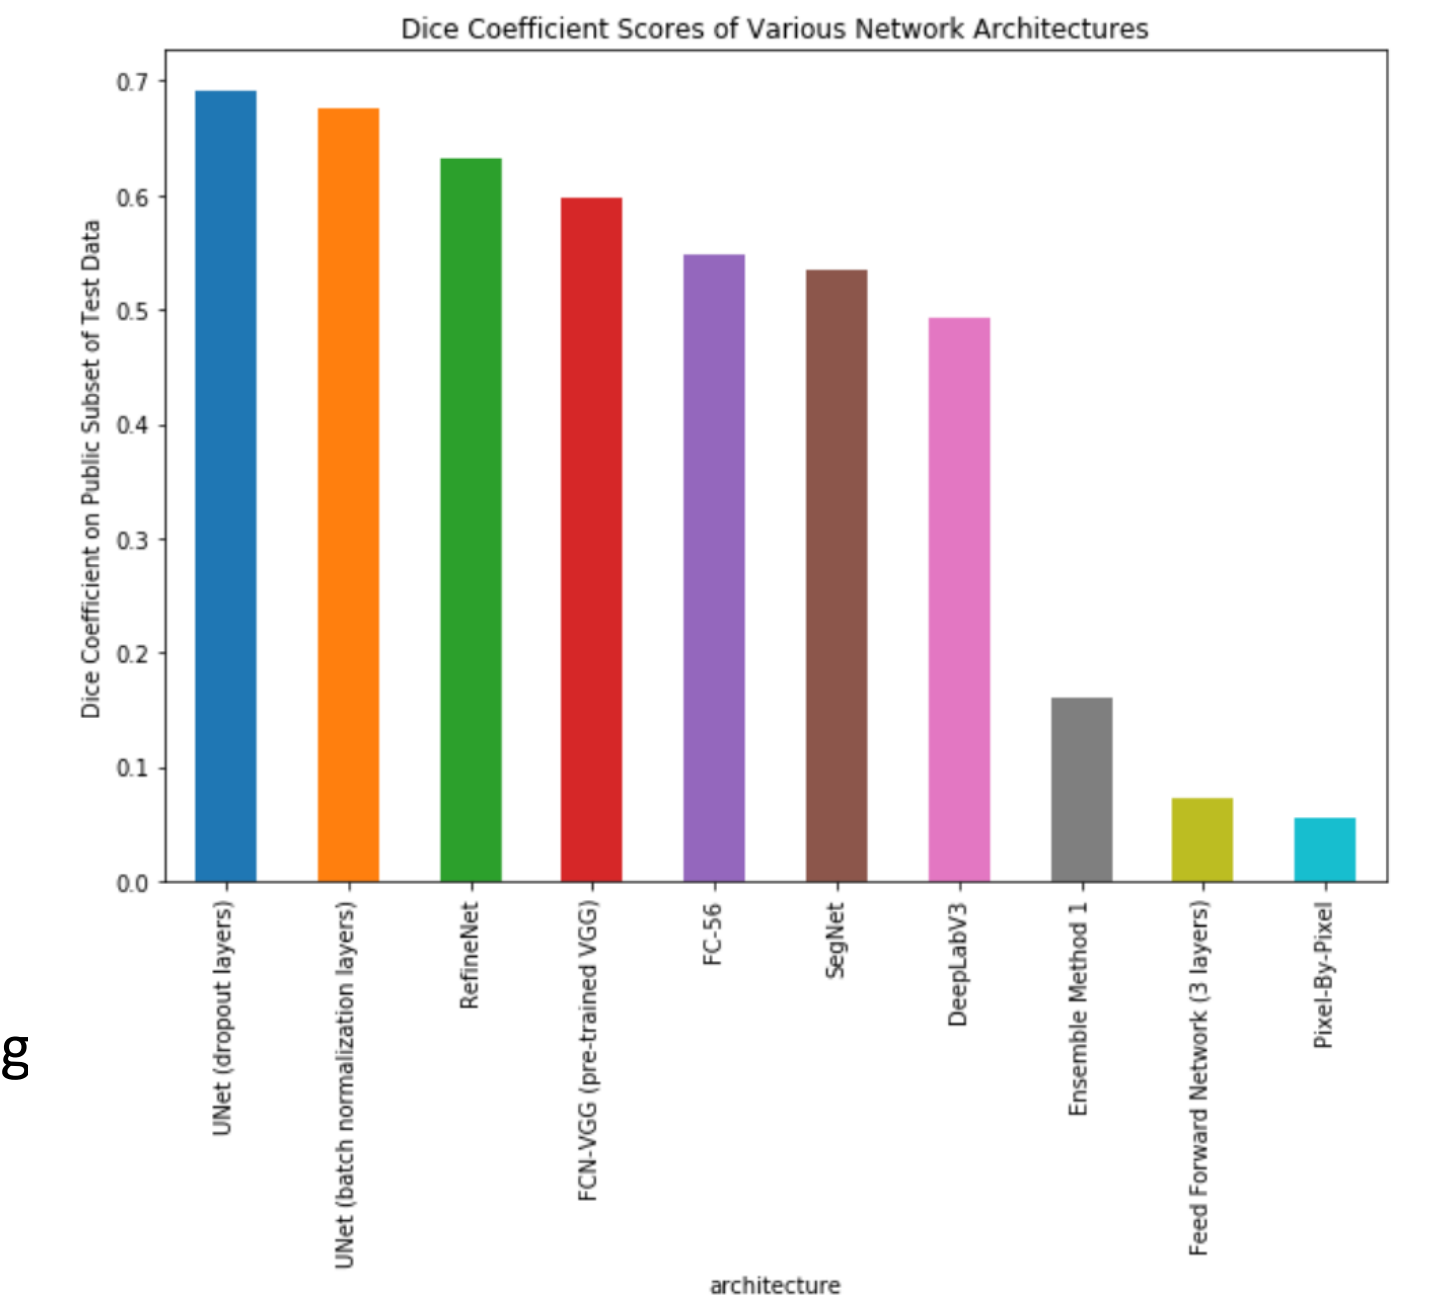


Training & Model Selection


- Models trained using Amazon Web Services & Google Cloud Platform
- Regularization using Dropout, and later Batch Normalization
- Standard Train-Validation split, evaluated after every epoch
- Scoring: Dice Coefficient

$$\frac{2TP}{2TP + FP + FN}$$

- Large dependency on recognizing road pixels



Timeline of Progress



- Massive spike when moving to Deep Convolutional Neural Networks (Late March)
- Progress slowed after tuning parameters for UNet
- Ensemble methods using majority vote and average value failed to improve

Future Directions

- Unsupervised Learning
 - Attempt to cluster types of images (country, city, etc) and build one model for each.
 - To predict on unseen images, classify based on clusters and then predict using relevant model
- Fix postprocessing using Conditional Random Fields
 - Road boundaries are almost always very smooth
 - Smooth our predictions after the fact to reflect this
- Integrated Conditional Random Fields within Neural Network Architecture
 - Instead of postprocessing, using a conditional random field as part of the network
 - Implement CRF as Recurrent Neural Network

Acknowledgements

Technical References
<https://arxiv.org/abs/1703.06870>
<https://devblogs.nvidia.com/solving-spacenet-road-detection-challenge-deep-learning/>
http://openaccess.thecvf.com/content_cvpr_2018_workshops/papers/w4/Buslaev_Fully_Convolutional_Network_CVPR_2018_paper.pdf
<https://medium.com/yemidalabs-innovation/data-augmentation-techniques-in-cnn-using-tensorflow-371ae43d5be9>
<https://medium.com/@rogerxuliang/setting-up-a-gpu-instance-for-deep-learning-on-aws-795343e16e44>
http://www.cs.toronto.edu/~fritz/absps/road_detection.pdf
<https://towardsdatascience.com/dont-use-dropout-in-convolutional-networks-81486c823c16>
<http://warmspringwinds.github.io/tensorflow/tf-slim/2016/12/18/image-segmentation-with-tensorflow-using-cnns-and-conditional-random-fields/>
<https://openreview.net/forum?id=B1Yy1Bx7Z>
Software
<https://github.com/GeorgeSeif/Semantic-Segmentation-Suite>
https://github.com/matterport/Mask_RCNN
https://github.com/sadeepi/crfasmn_keras
<https://github.com/lucasb-eyer/pydensecrf>
<https://faiyoonice.github.io/Learn-about-Fully-Convolutional-Networks-for-semantic-segmentation.html>
<https://www.kaggle.com/keegil/keras-u-net-starter-lb-0-277>